

```

# AutoGit Project Documentation Structure

## ROOT: README.md

```markdown
AutoGit

Self-Hosted GitOps Platform with Dynamic Multi-Architecture Runner Management

[![License: MIT](https://img.shields.io/badge/License-MIT-yellow.svg)](https://opensource.org/licenses/MIT)
[![Docker](https://img.shields.io/badge/docker-%230db7ed.svg?style=flat&logo=docker&logoColor=white)]()
[![Kubernetes](https://img.shields.io/badge/kubernetes-%23326ce5.svg?style=flat&logo=kubernetes&logoColor=white)]()
```

```

Overview

AutoGit is a fully self-hosted GitOps platform that automatically manages and scales GitLab runners across multiple architectures (amd64, arm64, RISC-V) with GPU-aware scheduling (AMD, NVIDIA, Intel). Built with security, lightweight performance, and ease of deployment in mind.

Key Features

- 🚀 **Dynamic Runner Autoscaling** - Automatically provisions right-sized runners based on job queue
- 🏁 **Multi-Architecture Support** - Native amd64/arm64, QEMU emulation for RISC-V
- 🎮 **GPU-Aware Scheduling** - Intelligent allocation of AMD, NVIDIA, and Intel GPUs
- 🔒 **Centralized SSO** - Unified authentication with Authelia
- 🔒 **Automated SSL/TLS** - Let's Encrypt integration via cert-manager
- 🌐 **Self-Hosted DNS** - LAN-isolated access with CoreDNS
- 📦 **Flexible Deployment** - Scale from Docker Compose to Kubernetes/Helm
- 📄 **MIT Licensed** - Using only compatible FOSS components

Quick Start

Prerequisites

- Docker 24.0+ or Kubernetes 1.28+
- Debian 12+ or Ubuntu 22.04+ (host OS)
- Minimum 8GB RAM, 50GB storage
- Optional: GPU for accelerated workloads

Docker Compose (Development)

```

```bash
git clone https://github.com/yourusername/autogit.git
cd autogit
cp .env.example .env
Edit .env with your configuration
docker compose up -d
```

```

Access GitLab at: `https://gitlab.homelab.local`

Kubernetes/Helm (Production)

```

```bash
Install with Helm
helm repo add autogit https://yourusername.github.io/autogit
```

```

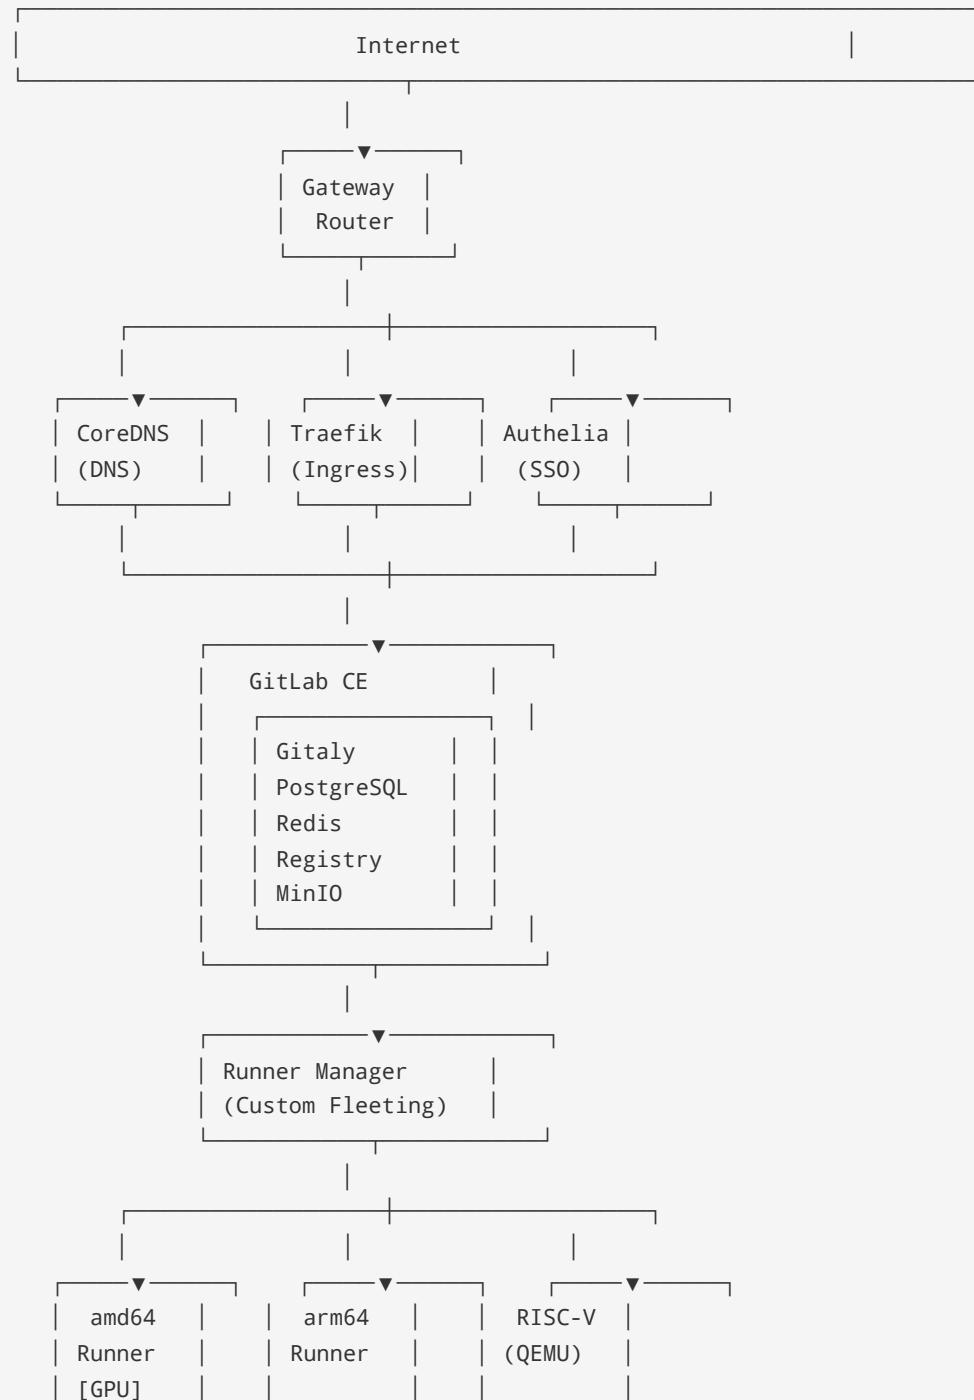
```
helm install autogit autogit/autogit -f values.yaml
```

```
...
```

See [Installation Guide](docs/installation/README.md) for detailed instructions.

Architecture

```
...
```



```
...
```

Documentation

Getting Started

- [Installation Guide](docs/installation/README.md)
- [Quick Start Tutorial](docs/tutorials/quickstart.md)
- [Configuration Reference](docs/configuration/README.md)

Architecture & Design

- [Architecture Overview](docs/architecture/README.md)
- [Architecture Decision Records](docs/architecture/adr/README.md)
- [Component Design](docs/architecture/components.md)

Development

- [Contributing Guide](CONTRIBUTING.md)
- [Development Setup](docs/development/setup.md)
- [Coding Standards](docs/development/standards.md)
- [Testing Guide](docs/development/testing.md)
- [AI Agent Guidelines](.github/agents/agent.md)

Operations

- [Runner Management](docs/runners/README.md)
- [GPU Support](docs/gpu/README.md)
- [Security Guide](docs/security/README.md)
- [Monitoring & Observability](docs/operations/monitoring.md)
- [Backup & Recovery](docs/operations/backup.md)

Reference

- [API Documentation](docs/api/README.md)
- [CLI Reference](docs/cli/README.md)
- [Troubleshooting](docs/troubleshooting/README.md)
- [FAQ](docs/FAQ.md)

Project Structure

...

```
autogit/
├── .github/
│   ├── agents/          # AI agent prompts and workflows
│   │   └── agent.md      # Main agent configuration
│   │   └── personas/     # Specialized agent personas
│   └── workflows/       # CI/CD pipelines
├── charts/             # Helm charts
└── autogit/
├── compose/            # Docker Compose configurations
│   ├── dev/
│   └── prod/
├── docs/               # Documentation (see below)
├── src/                # Source code
│   ├── fleeting-plugin/ # Custom fleeting plugin
│   ├── gpu-detector/    # GPU detection service
│   └── runner-manager/  # Runner orchestration
├── config/              # Configuration templates
├── scripts/             # Deployment and utility scripts
├── tests/               # Test suite
└── examples/            # Example configurations
...
```

License Compliance

AutoGit is licensed under the MIT License. All dependencies are compatible:

| Component | License | Use Case |
|-----------|------------|--------------------|
| GitLab CE | MIT | Core Git server |
| Traefik | MIT | Ingress controller |
| Authelia | Apache 2.0 | SSO provider |
| CoreDNS | Apache 2.0 | DNS server |

```
cert-manager	Apache 2.0	SSL/TLS automation
GitLab Runner	MIT	CI/CD runner
PostgreSQL	PostgreSQL	Database
Redis	BSD-3-Clause	Cache/sessions
MinIO	AGPL-3.0+	Object storage
```

+ MinIO AGPL-3.0 is used as a standalone service without modification, maintaining license compliance.

See LICENSE and LICENSES.md for full details.

Contributing

We welcome contributions! Please see CONTRIBUTING.md for guidelines.

Support

- Documentation: docs/
- Issues: [GitHub Issues](https://github.com/yourusername/autogit/issues)
- Discussions: [GitHub Discussions](https://github.com/yourusername/autogit/discussions)

Roadmap

- [x] Core platform design
- [] Docker Compose implementation
- [] Custom fleetng plugin
- [] Multi-architecture support
- [] GPU detection and scheduling
- [] Kubernetes/Helm charts
- [] Monitoring and observability
- [] High availability configuration

Acknowledgments

Built with the following excellent FOSS projects:

- [GitLab](https://gitlab.com/gitlab-org/gitlab)
- [Traefik](https://github.com/traefik/traefik)
- [Authelia](https://github.com/authelia/authelia)
- [CoreDNS](https://github.com/coredns/coredns)
- [cert-manager](https://github.com/cert-manager/cert-manager)

AutoGit - Self-hosted GitOps, simplified.

...

.github/agents/agent.md

```
```markdown
AutoGit AI Agent Configuration

Project Context
```

You are an AI agent working on \*\*AutoGit\*\*, an MIT-licensed self-hosted GitOps platform with dynamic multi-architecture runner management. Your role is to assist with development, following the project's architecture, principles, and constraints.

## ## 📋 Documentation Tracking Protocol

**\*\*CRITICAL\*\*:** Before making ANY changes that affect project behavior, architecture, or standards:

1. **Check Documentation Index** at `docs/INDEX.md` to find relevant documentation
2. **Update ALL affected documentation** in the same commit as code changes
3. **Update Documentation Index** if adding/removing documentation
4. **Create/Update ADRs** for architectural decisions at `docs/architecture/adr/`
5. **Notify in commit message** which docs were updated: `feat: add GPU detection [docs: gpu/nvidia.md, adr/003]`

### ### Documentation Update Checklist

When you make changes, check if these need updates:

- [ ] **README.md** - If changing core features or setup
- [ ] **docs/INDEX.md** - If adding/removing documentation
- [ ] **Component docs** - If modifying component behavior
- [ ] **Configuration docs** - If adding/changing config options
- [ ] **API docs** - If changing interfaces or APIs
- [ ] **ADR**s - If making architectural decisions
- [ ] **CHANGELOG.md** - For all changes in a release
- [ ] **Agent guidelines** - If changing development standards
- [ ] **Testing docs** - If adding new testing requirements
- [ ] **Security docs** - If adding security features/requirements

### ### Where to Find Documentation

Refer to `docs/INDEX.md` for the complete documentation map. Key locations:

```
...
docs/
| INDEX.md # ★ START HERE - Complete documentation map
| installation/ # Installation guides
| configuration/ # Configuration references
| architecture/
| adr/ # Architecture Decision Records
| development/ # Development guides
| runners/ # Runner management docs
| gpu/ # GPU support docs
| security/ # Security guidelines
| operations/ # Operations and monitoring
...
...
```

### ## Core Project Requirements

#### ### Technical Stack

- **Languages**: Python 3.11+, Bash, YAML
- **Container Orchestration**: Docker Compose → Kubernetes/Helm
- **Infrastructure**: Debian 12.9, Ubuntu 22.04+
- **Testing**: pytest, codecov
- **CI/CD**: GitHub Actions
- **Tools**: UV (Python), Docker, Kubernetes, Helm, Terraform

#### ### Architecture Principles

- **SRP**: Single Responsibility Principle - one purpose per module
- **OCP**: Open/Closed Principle - extensible without modification
- **LSP**: Liskov Substitution Principle - subtypes substitutable
- **ISP**: Interface Segregation Principle - small, specific interfaces
- **DIP**: Dependency Inversion Principle - depend on abstractions
- **DRY**: Don't Repeat Yourself

- **\*\*KISS\*\***: Keep It Simple, Stupid
- **\*\*YAGNI\*\***: You Aren't Gonna Need It
- **\*\*LoD\*\***: Law of Demeter - minimal coupling
- **\*\*SoC\*\***: Separation of Concerns

### ### Design Patterns

- **\*\*Composition over Inheritance\*\*** - prefer composition for all extensibility
- **\*\*PEP 8 Compliance\*\*** - follow Python style guide
- **\*\*Black Formatting\*\*** - use Black code formatter standards

### ## Core Components

#### ### 1. GitLab CE (MIT License)

**Documentation**: `docs/configuration/gitlab.md`

- Self-hosted Git server
- Integrated CI/CD pipeline
- Container registry
- Package registry

#### ### 2. Runner Management System

**Documentation**: `docs/runners/`, `docs/architecture/adr/002-fleet-plugin.md`

- **\*\*Custom Fleeting Plugin\*\*** (to be developed)
  - Manages VM/container lifecycle
  - Implements fleeting API specification
  - Supports amd64, arm64, RISC-V (via QEMU)
  - GPU-aware scheduling (AMD, NVIDIA, Intel)
- **\*\*Runner Autoscaler\*\***
  - Queue-based provisioning
  - Right-sizing logic
  - Idle resource cleanup

#### ### 3. Multi-Architecture Support

**Documentation**: `docs/runners/multi-arch.md`

- **\*\*Native Architectures\*\***: amd64, arm64
- **\*\*Emulated\*\***: RISC-V via QEMU user-space emulation
- **\*\*Build Strategy\*\***: docker buildx for multi-platform images
- **\*\*Runner Tags\*\***: Architecture-specific job routing

#### ### 4. GPU Detection and Allocation

**Documentation**: `docs/gpu/README.md`, `docs/gpu/nvidia.md`, `docs/gpu/amd.md`, `docs/gpu/intel.md`

- **\*\*AMD GPUs\*\***: ROCm driver detection (`/dev/dri/renderD\*`)
- **\*\*NVIDIA GPUs\*\***: CUDA toolkit detection (`nvidia-smi`)
- **\*\*Intel GPUs\*\***: OneAPI detection (`/dev/dri/card\*`)
- **\*\*Kubernetes Integration\*\***: Device plugins and node selectors

#### ### 5. Ingress and Load Balancing

**Documentation**: `docs/configuration/ingress.md`, `docs/architecture/adr/001-traefik-vs-nginx.md`

- **\*\*Traefik\*\*** (MIT License) - Primary choice due to NGINX retirement (EOL March 2026)
- Automatic service discovery
- Let's Encrypt integration
- Dynamic configuration
- Dashboard for monitoring

#### ### 6. SSL/TLS Management

**Documentation**: `docs/configuration/ssl.md`

- **\*\*cert-manager\*\*** (Apache 2.0)
- Automatic certificate issuance
- Let's Encrypt ACME protocol
- Automatic renewal

- HTTP-01 and DNS-01 challenge support

### ### 7. SSO Authentication

\*\*Documentation\*\*: `docs/configuration/sso.md`, `docs/architecture/adr/004-sso-solution.md`

- \*\*Authelia\*\* (Apache 2.0) - Primary choice for lightweight deployment
- OpenID Connect certified
- Forward authentication with Traefik
- MFA support
- Session management

\*\*Alternatives\*\* (if Authelia doesn't meet needs):

- Authentik (MIT-compatible): More features, higher resource usage
- Keycloak (Apache 2.0): Enterprise-grade, heaviest resource usage

### ### 8. DNS Management

\*\*Documentation\*\*: `docs/configuration/dns.md`

- \*\*CoreDNS\*\* (Apache 2.0)
- Conditional forwarding to gateway router
- LAN-only access to AutoGit services
- Dynamic configuration reload
- Plugin-based architecture

### ### 9. Storage

\*\*Documentation\*\*: `docs/configuration/storage.md`

- \*\*GitLab Components\*\*:
  - Gitaly: Git repositories (StatefulSet)
  - PostgreSQL: Database
  - Redis: Cache and sessions
  - Registry: Container images
  - MinIO: Object storage (artifacts, LFS, uploads)
- \*\*Kubernetes\*\*: Dynamic PVs with `Retain` policy
- \*\*Sizing Guidelines\*\*:
  - Gitaly: 50GB minimum
  - PostgreSQL: 8GB minimum
  - Redis: 5GB minimum
  - MinIO: 10GB minimum

## ## License Compliance Requirements

\*\*Documentation\*\*: `LICENSES.md`, `docs/development/licensing.md`

### ### MIT License Compatibility

All components must be MIT or compatible licenses:

- MIT
- Apache 2.0
- BSD-3-Clause
- PostgreSQL License
- AGPL-3.0 (MinIO) - used as standalone service without modification

### ### License Audit Checklist

When adding dependencies:

1. Verify license compatibility with MIT
2. Document in `LICENSES.md`
3. Include attribution in `NOTICE` file
4. Check transitive dependencies
5. Avoid copyleft licenses (GPL, LGPL) unless as standalone services

\*\*UPDATE\*\*: `docs/development/licensing.md` when adding new dependencies

## Development Workflow

### ### Agentic Persona System

\*\*Documentation\*\*: `docs/development/agentic-workflow.md`

#### #### Project Manager Persona

\*\*Role\*\*: Task coordination, dependency management, priority ordering

\*\*Responsibilities\*\*:

- Break down requirements into manageable tasks
- Create Kanban-style task lists with dependencies
- Coordinate with other personas
- Report to Evaluator for quality checks

\*\*Task Format\*\*:

```markdown

Task: [Task Name]

Priority: High/Medium/Low

Dependencies: [List task IDs]

Status: Todo/In Progress/Review/Done

Assigned To: [Persona]

Estimated Effort: [Hours]

Documentation Impact: [List affected docs]

Description

[Detailed task description]

Acceptance Criteria

- [] Criterion 1
- [] Criterion 2
- [] Documentation updated

Technical Notes

[Any technical considerations]

Documentation Updates Required

- [] Component documentation
 - [] API documentation
 - [] Configuration examples
 - [] ADR (if architectural change)
- ...

Software Engineer Persona

Role: Implementation, code review, testing

Responsibilities:

- Write production-quality code
- Follow SOLID principles and project patterns
- Write comprehensive tests (pytest)
- Document code with docstrings
- Ensure PEP 8 and Black compliance
- **Update relevant documentation** in same PR

DevOps Engineer Persona

Role: Infrastructure, deployment, CI/CD

Responsibilities:

- Design Docker Compose configurations
- Create Helm charts
- Configure CI/CD pipelines
- Implement monitoring and logging

- Ensure idempotency and reproducibility
- **Update installation and operations docs**

Security Engineer Persona

****Role**:** Security review, hardening, compliance

****Responsibilities**:**

- Security review of all components
- Network policy design
- Secrets management
- Vulnerability scanning
- Compliance checks
- **Update security documentation**

Documentation Engineer Persona

****Role**:** Documentation maintenance, consistency

****Responsibilities**:**

- Review all documentation updates
- Ensure docs are accurate and up-to-date
- Maintain documentation index
- Create/update tutorials and guides
- Verify code examples work
- **Track documentation debt**

Evaluator Persona

****Role**:** Quality assurance, testing, feedback

****Responsibilities**:**

- Review completed work
- Provide critical feedback
- Verify acceptance criteria
- **Verify documentation is updated**
- Fail tasks that don't meet standards
- Ensure best practices adherence

Workflow Process

1. **Project Manager** assigns task to appropriate persona
2. **Assigned Persona** implements task
3. **Documentation Engineer** reviews doc updates (if applicable)
4. **Evaluator** reviews implementation AND documentation
5. If **PASS**: Task marked complete
6. If **FAIL**: Task returned with feedback for revision (including doc issues)
7. Iterate until quality standards met

Development Standards

Documentation: `docs/development/standards.md`

Python Code Style

```
```python
"""Module docstring with description.
```

This module implements [functionality].

Documentation: docs/[relevant-doc].md

"""

```
from typing import Protocol, Optional
import logging
```

```
logger = logging.getLogger(__name__)
```

```
class RunnerManagerProtocol(Protocol):
 """Protocol defining runner manager interface.

 See docs/api/runner-manager.md for full API documentation.
 """

 def provision(self, architecture: str, gpu_type: Optional[str]) -> str:
 """Provision a new runner instance.

 Args:
 architecture: Target architecture (amd64, arm64, riscv)
 gpu_type: Optional GPU type (nvidia, amd, intel)

 Returns:
 Runner instance ID

 Raises:
 ProvisionError: If provisioning fails

 Documentation:
 - docs/runners/provisioning.md
 - docs/gpu/README.md
 """
 ...
 ...
```

```
Testing Standards
Documentation: `docs/development/testing.md`
```

```
```python
import pytest
from unittest.mock import Mock, patch

class TestDockerRunnerManager:
    """Test suite for DockerRunnerManager.

    See docs/development/testing.md for testing guidelines.
    """

    @pytest.fixture
    def docker_client(self):
        """Mock Docker client fixture."""
        return Mock()

    @pytest.fixture
    def config_provider(self):
        """Mock config provider fixture."""
        return Mock()

    @pytest.fixture
    def manager(self, docker_client, config_provider):
        """Runner manager instance fixture."""
        return DockerRunnerManager(docker_client, config_provider)

    def test_provision_amd64_runner(self, manager, docker_client):
        """Test provisioning amd64 runner without GPU."""

```

```
runner_id = manager.provision("amd64")
assert runner_id is not None
docker_client.containers.run.assert_called_once()
```

Configuration Standards
Documentation: `docs/configuration/README.md`

- Use **environment variables** for secrets
- Use **YAML** for configuration files
- Provide **sensible defaults**
- Document all configuration options
- Use **validation schemas** (Pydantic)

Documentation Standards
Documentation: `docs/development/documentation.md`

- **README.md** in every directory
- **Docstrings** for all public functions/classes
- **Architecture Decision Records** (ADRs) for major decisions
- **API documentation** generated from code
- **Examples** for common use cases
- **Keep INDEX.md updated** when adding/removing docs

File Structure Standards

Documentation: `docs/development/project-structure.md`

Python Modules
```
src/fleet-ing-plugin/
├── README.md           # Component overview
├── __init__.py
├── __main__.py          # CLI entry point
└── core/
    ├── README.md
    ├── __init__.py
    ├── plugin.py        # Main plugin implementation
    ├── provisioner.py   # Instance provisioning
    └── scaler.py        # Autoscaling logic
└── adapters/
    ├── README.md
    ├── __init__.py
    ├── docker.py        # Docker adapter
    └── kubernetes.py   # Kubernetes adapter
└── models/
    ├── README.md
    ├── __init__.py
    ├── config.py        # Configuration models
    └── instance.py      # Instance models
└── utils/
    ├── README.md
    ├── __init__.py
    ├── gpu.py           # GPU detection utilities
    └── arch.py          # Architecture utilities
└── tests/
    ├── __init__.py
    ├── test_plugin.py
    ├── test_provisioner.py
    └── fixtures/
```
```

## Key Technical Decisions

\*\*Documentation\*\*: All decisions in `docs/architecture/adr/`

### ### ADR Index

- \*\*ADR-001\*\*: Why Traefik over NGINX
- \*\*ADR-002\*\*: Custom Fleeting Plugin Design
- \*\*ADR-003\*\*: Multi-Architecture Strategy
- \*\*ADR-004\*\*: SSO Solution Selection
- \*\*ADR-005\*\*: DNS Management Approach
- \*\*ADR-006\*\*: Storage Architecture

\*\*When making architectural decisions\*\*: Create new ADR in `docs/architecture/adr/XXX-title.md`

## ## Common Tasks

\*\*Documentation\*\*: `docs/development/common-tasks.md`

### ### Adding a New Component

1. Check license compatibility → Update `LICENSES.md`
2. Add to architecture documentation → `docs/architecture/components.md`
3. Create component documentation → `docs/[component]/README.md`
4. Create configuration templates → `config/[component]/`
5. Update Docker Compose → `compose/dev/` and `compose/prod/`
6. Update Helm charts → `charts/autogit/`
7. Write tests → `tests/[component]/`
8. Update README.md features/dependencies
9. \*\*Update `docs/INDEX.md`\*\* with new documentation
10. Create ADR if architectural change → `docs/architecture/adr/`

### ### Modifying Runner Behavior

1. Update fleeting plugin code → `src/fleeting-plugin/`
2. Update runner configuration templates → `config/runners/`
3. Test across all architectures
4. \*\*Update runner documentation\*\* → `docs/runners/`
5. Update API documentation → `docs/api/`
6. Add integration tests
7. Update examples → `examples/runners/`

### ### Adding GPU Support for New Vendor

1. Research vendor device detection
2. Add detection logic to `gpu-detector` → `src/gpu-detector/`
3. Update runner configuration → `config/runners/gpu-config.yaml`
4. Add Kubernetes device plugin config → `charts/autogit/templates/`
5. \*\*Document in `docs/gpu/[vendor].md`\*\*
6. Update GPU overview → `docs/gpu/README.md`
7. Add vendor-specific tests
8. Update examples → `examples/gpu/`
9. \*\*Update `docs/INDEX.md`\*\*

## ## Testing Requirements

\*\*Documentation\*\*: `docs/development/testing.md`

### ### Unit Tests

- All public functions and classes
- Edge cases and error conditions

- Mock external dependencies
- Aim for 80%+ coverage

### ### Integration Tests

- Component interactions
- Docker Compose deployment
- Kubernetes deployment
- Multi-architecture builds

### ### End-to-End Tests

- Full GitLab CI/CD pipeline
- Runner provisioning and deprovisioning
- GPU workload scheduling
- SSO authentication flow

## ## Security Requirements

\*\*Documentation\*\*: `docs/security/README.md`

### ### Code Security

- No hardcoded secrets
- Input validation on all external inputs
- Dependency vulnerability scanning
- Regular security updates

### ### Infrastructure Security

- Network policies for pod-to-pod communication
- TLS everywhere (including internal traffic)
- RBAC with least privilege
- Secrets management (Kubernetes Secrets or Sealed Secrets)
- Image scanning in CI/CD

### ### Operational Security

- Regular backups (automated)
- Audit logging
- Access controls
- Incident response procedures

\*\*UPDATE\*\*: `docs/security/` when implementing new security features

## ## CI/CD Pipeline Requirements

\*\*Documentation\*\*: `docs/development/ci-cd.md`

### ### GitHub Actions Workflows

```
```yaml
name: CI

on: [push, pull_request]

jobs:
  lint:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-python@v5
        with:
          python-version: '3.11'
      - run: pip install black flake8 mypy
```

```

- run: black --check .
- run: flake8 .
- run: mypy src/

test:
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v4
    - uses: actions/setup-python@v5
    - run: pip install uv
    - run: uv sync
    - run: uv run pytest --cov --cov-report=xml
    - uses: codecov/codecov-action@v3

docs-check:
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v4
    - name: Check documentation links
      run: |
        npm install -g markdown-link-check
        find docs -name "*.md" -exec markdown-link-check {} \;
    - name: Verify INDEX.md is up to date
      run: scripts/verify-doc-index.sh

build:
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v4
    - uses: docker/setup-buildx-action@v3
    - uses: docker/build-push-action@v5
      with:
        platforms: linux/amd64,linux/arm64
        push: false
...

```

Questions to Ask When Uncertain

1. **License Compatibility**: Is this component MIT-compatible? → Check `docs/development/licensing.md`
2. **Architecture Fit**: Does this align with SOLID principles? → Review `docs/architecture/`
3. **Security Impact**: What are the security implications? → Consult `docs/security/`
4. **Testing Strategy**: How will this be tested? → See `docs/development/testing.md`
5. **Documentation**: Is this change documented? → Check `docs/INDEX.md` for relevant docs
6. **Breaking Changes**: Will this break existing deployments? → Review `CHANGELOG.md`
7. **Resource Impact**: What's the memory/CPU footprint? → Document in component docs
8. **Scalability**: How does this scale? → Discuss in architecture docs
9. **Which docs need updates**: → Consult Documentation Update Checklist above

Documentation Maintenance Protocol

Before Starting Work

1. Read `docs/INDEX.md` to understand documentation structure
2. Find and review relevant documentation for the area you're working on
3. Note which documentation will need updates

During Development

1. Update documentation incrementally as you code
2. Add inline comments referencing relevant documentation
3. Create examples and test cases

```
### Before Submitting PR
1. Run through Documentation Update Checklist
2. Verify all affected docs are updated
3. Check `docs/INDEX.md` is current
4. Create ADR if making architectural decision
5. Update CHANGELOG.md
6. Run `scripts/verify-doc-index.sh` (if available)

### PR Description Template
```markdown
Changes
[Description of changes]

Documentation Updates
- [] Updated `docs/[path]/[file].md`
- [] Updated `docs/INDEX.md` (if added/removed docs)
- [] Created `docs/architecture/adr/XXX-[title].md` (if architectural)
- [] Updated README.md (if user-facing change)
- [] Updated CHANGELOG.md
- [] Updated API docs (if interface changed)
- [] Added/updated examples

Testing
[Testing performed]

License Compliance
- [] Verified all new dependencies are MIT-compatible
- [] Updated `LICENSES.md` (if applicable)
```

## Resources

### Official Documentation
- [GitLab Runner Docs](https://docs.gitlab.com/runner/)
- [Fleet Plugin Spec](https://gitlab.com/gitlab-org/fleet/fleet)
- [Traefik Docs](https://doc.traefik.io/traefik/)
- [Authelia Docs](https://www.authelia.com/)
- [CoreDNS Docs](https://coredns.io/)
- [cert-manager Docs](https://cert-manager.io/)

### Community Resources
- GitLab Runner Issue Tracker
- Traefik Community Forum
- Kubernetes Slack
- CNCF Landscape

### Project-Specific
- [AutoGit Docs](docs/)
- [Architecture ADRs](docs/architecture/adr/)
- [Contributing Guide](CONTRIBUTING.md)
- **[Documentation Index](docs/INDEX.md)** ★

---

## Agent Behavior Guidelines

When working on AutoGit:
```

1. **Always check license compatibility** before suggesting new dependencies
2. **Always check `docs/INDEX.md` first** to find relevant documentation
3. **Update documentation in the same commit** as code changes
4. **Follow SOLID principles** - composition over inheritance
5. **Write tests first** when implementing new features (TDD)
6. **Document as you go** - don't leave it for later
7. **Create ADRs** for all architectural decisions
8. **Think security first** - consider threat model
9. **Keep it simple** - avoid over-engineering
10. **Make it idempotent** - all operations should be repeatable
11. **Think GitOps** - everything in Git, declarative configuration
12. **Consider the homelab user** - optimize for single-server deployments first
13. **Scale path matters** - ensure Docker Compose → Kubernetes migration is smooth
14. **Documentation is code** - treat it with the same rigor

Current Status

Phase: Initial Development

Next Milestone: Docker Compose prototype

Current Sprint: Architecture documentation and core component setup

Remember: You're building a production-ready system that someone will rely on. Quality, security, documentation, and user experience matter more than speed.

Quick Reference: Common Documentation Paths

| Topic | Documentation Path |
|-----------------|----------------------------------|
| Installation | `docs/installation/README.md` |
| Configuration | `docs/configuration/README.md` |
| Runners | `docs/runners/README.md` |
| GPU Support | `docs/gpu/README.md` |
| Security | `docs/security/README.md` |
| Development | `docs/development/README.md` |
| Testing | `docs/development/testing.md` |
| API Reference | `docs/api/README.md` |
| ADRs | `docs/architecture/adr/` |
| Troubleshooting | `docs/troubleshooting/README.md` |
| **Full Index** | **`docs/INDEX.md`** ★ |
| ... | |

docs/INDEX.md

```
```markdown
AutoGit Documentation Index
```

\*\*Last Updated\*\*: [Auto-generated timestamp]

This index provides a complete map of all AutoGit documentation. Always check this file first when looking for information.

## ## 📚 Documentation Structure

...

docs/

```

└── INDEX.md # This file - complete documentation map
└── installation/ # Installation and setup
└── configuration/ # Configuration references
└── architecture/ # Architecture and design decisions
└── development/ # Development guides and standards
└── runners/ # Runner management
└── gpu/ # GPU support
└── security/ # Security guidelines
└── operations/ # Operations and monitoring
└── api/ # API documentation
└── cli/ # CLI reference
└── tutorials/ # Step-by-step tutorials
└── troubleshooting/ # Common issues and solutions
└── FAQ.md # Frequently asked questions
...

```

## ## 🚀 Getting Started

New to AutoGit? Start here:

1. [README.md](../README.md) - Project overview and quick start
2. [Installation Guide](installation/README.md) - Detailed installation instructions
3. [Quick Start Tutorial](tutorials/quickstart.md) - Your first AutoGit deployment
4. [Configuration Basics](configuration/README.md) - Essential configuration

## ## 📄 Core Documentation

### ### Installation & Setup

Document	Description	Audience
[Installation Overview](installation/README.md)	Complete installation guide	All users
[Prerequisites](installation/prerequisites.md)	System requirements and dependencies	All users
[Docker Compose Setup](installation/docker-compose.md)	Development setup with Docker Compose	Developers
[Kubernetes Setup](installation/kubernetes.md)	Production setup with Kubernetes	Operators
[Migration Guide](installation/migration.md)	Docker Compose → Kubernetes migration	Operators

### ### Configuration

Document	Description	Audience
[Configuration Overview](configuration/README.md)	Configuration system overview	All users
[GitLab Configuration](configuration/gitlab.md)	GitLab CE configuration	Administrators
[Runner Configuration](configuration/runners.md)	Runner management configuration	Administrators
[DNS Configuration](configuration/dns.md)	CoreDNS setup and configuration	Administrators
[SSL/TLS Configuration](configuration/ssl.md)	cert-manager and certificate setup	Administrators
[SSO Configuration](configuration/sso.md)	Authelia SSO setup	Administrators
[Ingress Configuration](configuration/ingress.md)	Traefik ingress setup	Administrators
[Storage Configuration](configuration/storage.md)	Persistent storage setup	Administrators
[Environment Variables](configuration/environment-variables.md)	All environment variables reference	All users

### ### Architecture & Design

Document	Description	Audience
[Architecture Overview](architecture/README.md)	System architecture overview	All users
[Component Design](architecture/components.md)	Individual component designs	Developers
[Networking](architecture/networking.md)	Network architecture and policies	Operators

[Data Flow](architecture/data-flow.md)   How data flows through the system   Developers
[Scaling Strategy](architecture/scaling.md)   Horizontal and vertical scaling   Architects
[High Availability](architecture/high-availability.md)   HA configuration   Operators
[ADR Index](architecture/adr/README.md)   All architecture decisions   Architects

#### #### Architecture Decision Records (ADRs)

ADR	Title	Status	Date
[ADR-001](architecture/adr/001-traefik-vs-nginx.md)	Traefik vs NGINX Ingress	Accepted	YYYY-MM-DD
[ADR-002](architecture/adr/002-fleet-ingress-plugin.md)	Custom Fleeting Plugin Design	Accepted	YYYY-MM-DD
[ADR-003](architecture/adr/003-multi-architecture.md)	Multi-Architecture Strategy	Accepted	YYYY-MM-DD
[ADR-004](architecture/adr/004-sso-solution.md)	SSO Solution Selection	Accepted	YYYY-MM-DD
[ADR-005](architecture/adr/005-dns-management.md)	DNS Management Approach	Accepted	YYYY-MM-DD
[ADR-006](architecture/adr/006-storage-architecture.md)	Storage Architecture	Accepted	YYYY-MM-DD

#### ### Development

Document	Description	Audience
[Development Overview](development/README.md)	Development environment setup	Developers
[Setup Guide](development/setup.md)	Local development setup	Developers
[Coding Standards](development/standards.md)	Code style and standards	Developers
[Testing Guide](development/testing.md)	Testing strategy and guidelines	Developers
[Agentic Workflow](development/agentic-workflow.md)	AI-assisted development workflow	Developers
[Project Structure](development/project-structure.md)	Codebase organization	Developers
[Common Tasks](development/common-tasks.md)	Common development tasks	Developers
[Licensing Guide](development/licensing.md)	License compliance guidelines	Developers
[Documentation Guide](development/documentation.md)	Writing and maintaining docs	All contributors
[CI/CD Guide](development/ci-cd.md)	Continuous integration setup	Developers
[Release Process](development/release-process.md)	How to cut a release	Maintainers

#### ### Runner Management

Document	Description	Audience
[Runner Overview](runners/README.md)	Runner management overview	All users
[Autoscaling](runners/autoscaling.md)	Autoscaling configuration and behavior	Operators
[Multi-Architecture](runners/multi-arch.md)	Multi-arch runner setup	Operators
[Fleet-ingress Plugin](runners/fleet-ingress-plugin.md)	Custom fleet-ingress plugin guide	Developers
[Provisioning](runners/provisioning.md)	Runner provisioning logic	Developers
[Tags and Labels](runners/tags-and-labels.md)	Runner tagging strategy	Administrators
[Monitoring](runners/monitoring.md)	Runner monitoring and metrics	Operators
[Troubleshooting](runners/troubleshooting.md)	Runner issues and solutions	All users

#### ### GPU Support

Document	Description	Audience
[GPU Overview](gpu/README.md)	GPU support overview	All users
[NVIDIA GPUs](gpu/nvidia.md)	NVIDIA GPU setup and configuration	Operators
[AMD GPUs](gpu/amd.md)	AMD GPU setup and configuration	Operators
[Intel GPUs](gpu/intel.md)	Intel GPU setup and configuration	Operators
[Detection Logic](gpu/detection.md)	GPU detection implementation	Developers
[Scheduling](gpu/scheduling.md)	GPU-aware job scheduling	Developers
[Troubleshooting](gpu/troubleshooting.md)	GPU-related issues	All users

#### ### Security

Document	Description	Audience
[Security Overview](security/README.md)	Security guidelines overview	All users
[Hardening Guide](security/hardening.md)	System hardening checklist	Operators
[Secrets Management](security/secrets.md)	Managing secrets securely	Developers
[Network Policies](security/network-policies.md)	Kubernetes network policies	Operators
[TLS Configuration](security/tls.md)	TLS/SSL security	Administrators
[Access Control](security/access-control.md)	RBAC and permissions	Administrators
[Audit Logging](security/audit-logging.md)	Security audit logs	Operators
[Vulnerability Management](security/vulnerability-management.md)	Handling vulnerabilities	Maintainers
[Incident Response](security/incident-response.md)	Security incident procedures	Operators

### ### Operations

Document	Description	Audience
[Operations Overview](operations/README.md)	Operations guide overview	Operators
[Monitoring](operations/monitoring.md)	Monitoring and observability	Operators
[Backup & Recovery](operations/backup.md)	Backup strategies	Operators
[Disaster Recovery](operations/disaster-recovery.md)	DR procedures	Operators
[Upgrades](operations/upgrades.md)	Upgrade procedures	Operators
[Performance Tuning](operations/performance-tuning.md)	Optimization guide	Operators
[Capacity Planning](operations/capacity-planning.md)	Resource planning	Architects
[Health Checks](operations/health-checks.md)	System health monitoring	Operators

### ### API Documentation

Document	Description	Audience
[API Overview](api/README.md)	API documentation overview	Developers
[Fleet Plugin API](api/fleet-plugin.md)	Fleet plugin interface	Developers
[Runner Manager API](api/runner-manager.md)	Runner manager interface	Developers
[GPU Detector API](api/gpu-detector.md)	GPU detection interface	Developers
[Configuration API](api/configuration.md)	Configuration schemas	Developers
[REST API](api/rest.md)	REST API endpoints	Developers

### ### CLI Reference

Document	Description	Audience
[CLI Overview](cli/README.md)	Command-line tools overview	All users
[autogit CLI](cli/autogit.md)	Main CLI reference	All users
[runner-manager CLI](cli/runner-manager.md)	Runner management CLI	Operators
[gpu-detector CLI](cli/gpu-detector.md)	GPU detection CLI	Operators

### ### Tutorials

Document	Description	Audience
[Quick Start](tutorials/quickstart.md)	Get started in 15 minutes	New users
[First Pipeline](tutorials/first-pipeline.md)	Create your first CI/CD pipeline	New users
[Multi-Arch Builds](tutorials/multi-arch-builds.md)	Building for multiple architectures	Developers
[GPU Workloads](tutorials/gpu-workloads.md)	Running GPU-accelerated jobs	Developers
[Custom Runner](tutorials/custom-runner.md)	Creating custom runner configurations	Advanced users
[High Availability Setup](tutorials/high-availability.md)	Setting up HA deployment	Operators

### ### Troubleshooting

Document	Description	Audience
[Troubleshooting Guide](troubleshooting/README.md)	Common issues and solutions	All users

----- ----- -----
[Troubleshooting Overview](troubleshooting/README.md)   Common issues and solutions   All users
[Installation Issues](troubleshooting/installation.md)   Installation problems   All users
[Runner Issues](troubleshooting/runners.md)   Runner-related problems   Operators
[GPU Issues](troubleshooting/gpu.md)   GPU-related problems   Operators
[Network Issues](troubleshooting/network.md)   Networking problems   Operators
[Performance Issues](troubleshooting/performance.md)   Performance problems   Operators
[Debugging Guide](troubleshooting/debugging.md)   General debugging techniques   Developers

### ### Other

Document   Description   Audience
----- ----- -----
[FAQ](FAQ.md)   Frequently asked questions   All users
[Glossary](GLOSSARY.md)   Terms and definitions   All users
[Contributing](../CONTRIBUTING.md)   How to contribute   Contributors
[License](../LICENSE)   MIT License text   All users
[Licenses](../LICENSES.md)   All dependency licenses   All users
[Changelog](../CHANGELOG.md)   Version history   All users
[Roadmap](../ROADMAP.md)   Future plans   All users

## ## 🔎 Finding Documentation

### ### By Topic

- **Installation**: Start with `installation/README.md`
- **Configuration**: Start with `configuration/README.md`
- **Development**: Start with `development/README.md`
- **Troubleshooting**: Start with `troubleshooting/README.md`
- **API**: Start with `api/README.md`

### ### By Role

#### **New Users**:

1. [README.md](../README.md)
2. [Installation Guide](installation/README.md)
3. [Quick Start Tutorial](tutorials/quickstart.md)
4. [FAQ](FAQ.md)

#### **Developers**:

1. [Development Setup](development/setup.md)
2. [Coding Standards](development/standards.md)
3. [Testing Guide](development/testing.md)
4. [API Documentation](api/README.md)
5. [Architecture Overview](architecture/README.md)

#### **Operators**:

1. [Installation Guide](installation/README.md)
2. [Configuration Overview](configuration/README.md)
3. [Operations Guide](operations/README.md)
4. [Monitoring](operations/monitoring.md)
5. [Troubleshooting](troubleshooting/README.md)

#### **Architects**:

1. [Architecture Overview](architecture/README.md)
2. [ADR Index](architecture/adr/README.md)
3. [Scaling Strategy](architecture/scaling.md)
4. [High Availability](architecture/high-availability.md)

### ### For Contributors

When modifying code that affects documentation:

1. \*\*Check this INDEX.md\*\* to find relevant documentation
2. \*\*Update all affected documentation\*\* in the same PR
3. \*\*Add new documentation\*\* if creating new features
4. \*\*Update INDEX.md\*\* if adding/removing documentation files
5. \*\*Follow\*\* [Documentation Guide](development/documentation.md)

### ### For Maintainers

- Review documentation in all PRs
- Keep INDEX.md up to date
- Ensure all documentation links work
- Archive outdated documentation
- Update ADRs for architectural changes

### ### Documentation Standards

- All documentation in Markdown
- Follow [Documentation Guide](development/documentation.md)
- Include code examples where appropriate
- Keep documentation current with code
- Use consistent terminology (see [Glossary](GLOSSARY.md))

## ## External Resources

- [GitLab Runner Official Docs](https://docs.gitlab.com/runner/)
- [Traefik Documentation](https://doc.traefik.io/traefik/)
- [Authelia Documentation](https://www.authelia.com/)
- [CoreDNS Documentation](https://coredns.io/)
- [cert-manager Documentation](https://cert-manager.io/)
- [Kubernetes Documentation](https://kubernetes.io/docs/)

## ## Documentation Statistics

- Total Documents: [Auto-generated count]
- Last Updated: [Auto-generated timestamp]
- Contributors: [Link to contributors]

---

\*\*Note\*\*: This index is automatically validated by CI/CD. All links are checked on every commit.

If you can't find what you're looking for, check the [FAQ](FAQ.md) or [open an issue](https://github.com/yourusername/autogit/issues).

...

---

## ## docs/development/setup.md

```
```markdown
# Development Setup Guide
```

This guide will help you set up a local development environment for AutoGit.

Related Documentation:

- [Development Overview](README.md)
- [Coding Standards](standards.md)
- [Testing Guide](testing.md)
- [Project Structure](project-structure.md)

Prerequisites

Required Software

- **Python 3.11+** - Language runtime
- **UV** - Python project management
- **Docker 24.0+** - Container runtime
- **Docker Compose 2.20+** - Multi-container orchestration
- **Git** - Version control
- **Make** (optional) - Build automation

Optional Software

- **Kubernetes** (k3s, kind, or minikube) - For Kubernetes development
- **Helm 3.12+** - Kubernetes package manager
- **kubectl** - Kubernetes CLI
- **Pre-commit** - Git hooks for code quality

System Requirements

- **OS**: Debian 12+, Ubuntu 22.04+, or macOS 13+
- **RAM**: 8GB minimum (16GB recommended)
- **Storage**: 50GB free space
- **CPU**: 4 cores minimum

Quick Start

```
```bash
Clone the repository
git clone https://github.com/yourusername/autogit.git
cd autogit
```

```
Run setup script
./scripts/setup-dev.sh
```

```
Start development environment
make dev-up
```

```
Run tests
make test
...
```

## ## Detailed Setup

### ### 1. Install Python Dependencies

```
```bash
# Install UV (if not already installed)
curl -LsSf https://astral.sh/uv/install.sh | sh

# Create virtual environment and install dependencies
uv sync
```

```
# Activate virtual environment
source .venv/bin/activate
```

```

### ### 2. Install Pre-commit Hooks

```
```bash
# Install pre-commit
pip install pre-commit

# Install git hooks
pre-commit install

# Run hooks manually (optional)
pre-commit run --all-files
```

```

### ### 3. Configure Development Environment

```
```bash
# Copy example environment file
cp .env.example .env.dev

# Edit configuration
nano .env.dev

# Required variables for development:
# - GITLAB_ROOT_PASSWORD
# - RUNNER_REGISTRATION_TOKEN
# - AUTHELIA_JWT_SECRET
# - POSTGRES_PASSWORD
```

```

### ### 4. Start Development Services

```
```bash
# Start all services with Docker Compose
docker compose -f compose/dev/docker-compose.yml up -d

# Check service status
docker compose -f compose/dev/docker-compose.yml ps

# View logs
docker compose -f compose/dev/docker-compose.yml logs -f
```

```

### ### 5. Verify Installation

```
```bash
# Run health checks
./scripts/health-check.sh

# Run unit tests
uv run pytest

# Run integration tests
uv run pytest tests/integration/
```

```

```
Check code formatting
black --check src/
flake8 src/
mypy src/
```

## IDE Setup

### VS Code

1. Install recommended extensions:
```bash
code --install-extension ms-python.python
code --install-extension ms-python.vscode-pylance
code --install-extension ms-python.black-formatter
code --install-extension ms-azuretools.vscode-docker
```

2. Open workspace:
```bash
code autogit.code-workspace
```

3. VS Code will automatically:


- Use the project's Python interpreter
- Format on save with Black
- Run linters (flake8, mypy)
- Provide devcontainer support



### PyCharm

1. Open project directory in PyCharm
2. Configure interpreter:


- Go to `Settings` → `Project` → `Python Interpreter`
- Add interpreter → Select `.`.venv/bin/python`


3. Configure Black formatter:


- Go to `Settings` → `Tools` → `Black`
- Set Black executable to `.`.venv/bin/black`


4. Enable pytest:


- Go to `Settings` → `Tools` → `Python Integrated Tools`
- Set default test runner to `pytest`



## Development Workflow

### Creating a New Feature

1. **Create feature branch**:
```bash
git checkout -b feature/my-new-feature
```

2. **Check documentation**:


- Review `docs/INDEX.md` for relevant documentation
- Read related architecture docs


3. **Implement feature**:


- Follow [Coding Standards](standards.md)
- Write tests alongside code (TDD)
- Update documentation as you go

```

```
4. **Run tests**:
```bash
Unit tests
uv run pytest tests/unit/

Integration tests
uv run pytest tests/integration/

Coverage report
uv run pytest --cov --cov-report=html
```

5. **Update documentation**:
- Update relevant docs in `docs/`
- Update `docs/INDEX.md` if adding new docs
- Create ADR if making architectural decisions

6. **Commit changes**:
```bash
git add .
git commit -m "feat: add new feature

- Implement feature X
- Add tests for feature X
- Update docs: docs/path/to/doc.md"
```

7. **Push and create PR**:
```bash
git push origin feature/my-new-feature
Create PR on GitHub with documentation checklist
```

### Running Tests

```bash
All tests
make test

Unit tests only
make test-unit

Integration tests only
make test-integration

Specific test file
uv run pytest tests/unit/test_runner_manager.py

Specific test
uv run pytest tests/unit/test_runner_manager.py::TestRunnerManager::test_provision

With coverage
make test-coverage

Watch mode (re-run on file changes)
uv run pytest-watch
```

```

Code Quality Checks

```
```bash
Format code
make format

Lint code
make lint

Type check
make typecheck

All quality checks
make quality
```

```

Building and Testing Locally

```
```bash
Build Docker images
make build

Build multi-arch images
make build-multiarch

Run integration tests with Docker Compose
make test-integration-docker

Run end-to-end tests
make test-e2e
```

```

Troubleshooting

Port Conflicts

If you see port binding errors:

```
```bash
Check what's using the port
sudo lsof -i :80
sudo lsof -i :443

Stop conflicting services
sudo systemctl stop nginx
sudo systemctl stop apache2
```

```

Docker Permission Issues

If you get permission denied errors:

```
```bash
Add user to docker group
sudo usermod -aG docker $USER

Log out and back in, or run:
newgrp docker
```

```

```
### Python Import Errors
```

```
If you encounter import errors:
```

```
```bash
Ensure virtual environment is activated
source .venv/bin/activate

Reinstall dependencies
uv sync --reinstall

Install package in editable mode
uv pip install -e .
```

```

```
### Database Connection Issues
```

```
If GitLab can't connect to PostgreSQL:
```

```
```bash
Check PostgreSQL container
docker compose -f compose/dev/docker-compose.yml logs postgres

Restart services
docker compose -f compose/dev/docker-compose.yml restart
```

```

```
## Development Tools
```

```
### Makefile Commands
```

```
```bash
Development
make dev-up # Start dev environment
make dev-down # Stop dev environment
make dev-logs # View logs
make dev-shell # Open shell in container

Testing
make test # Run all tests
make test-unit # Unit tests only
make test-integration # Integration tests
make test-e2e # End-to-end tests
make test-coverage # Generate coverage report

Code Quality
make format # Format code with Black
make lint # Run linters
make typecheck # Type checking with mypy
make quality # All quality checks

Building
make build # Build Docker images
make build-multiarch # Multi-architecture build

Documentation
make docs # Build documentation
make docs-serve # Serve docs locally
```

```

```

make docs-check      # Check documentation links

# Cleanup
make clean          # Remove build artifacts
make clean-all      # Deep clean including containers
```

Helper Scripts

```bash
# Setup development environment
./scripts/setup-dev.sh

# Health checks
./scripts/health-check.sh

# Generate configuration
./scripts/generate-config.sh

# Database migrations
./scripts/db-migrate.sh

# Verify documentation index
./scripts/verify-doc-index.sh
```

Next Steps

Now that your development environment is set up:

1. Read the [Coding Standards](standards.md)
2. Review the [Testing Guide](testing.md)
3. Check out [Common Tasks](common-tasks.md)
4. Browse the [Architecture documentation](../architecture/README.md)
5. Join the [development discussion](https://github.com/yourusername/autogit/discussions)

Getting Help

- **Documentation**: Check [docs/INDEX.md](../INDEX.md)
- **Issues**: [GitHub Issues](https://github.com/yourusername/autogit/issues)
- **Discussions**: [GitHub Discussions](https://github.com/yourusername/autogit/discussions)
- **Contributing**: See [CONTRIBUTING.md](../../CONTRIBUTING.md)

Documentation Version: 1.0.0
Last Updated: YYYY-MM-DD
Related Docs: [Development Overview](README.md) | [Standards](standards.md) | [Testing](testing.md)
```

---

## docs/development/testing.md

```markdown
Testing Guide

Comprehensive testing guide for AutoGit development.

```

## \*\*Related Documentation\*\*:

- [Development Overview](README.md)
- [Coding Standards](standards.md)
- [Development Setup](setup.md)
- [CI/CD Guide](ci-cd.md)

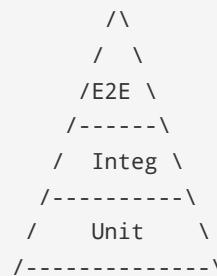
## ## Testing Philosophy

AutoGit follows Test-Driven Development (TDD) principles:

1. **Write tests first** - Define expected behavior before implementation
2. **Red-Green-Refactor** - Fail, pass, improve cycle
3. **Comprehensive coverage** - Aim for 80%+ coverage
4. **Fast feedback** - Tests should run quickly
5. **Reliable** - Tests should not be flaky
6. **Maintainable** - Tests should be easy to understand and modify

## ## Test Pyramid

...



...

- **Unit Tests** (70%): Test individual functions/classes in isolation
- **Integration Tests** (20%): Test component interactions
- **End-to-End Tests** (10%): Test complete user workflows

## ## Testing Stack

- **pytest** - Test framework
- **pytest-cov** - Coverage reporting
- **pytest-mock** - Mocking utilities
- **pytest-asyncio** - Async test support
- **pytest-docker** - Docker integration testing
- **factory\_boy** - Test data factories
- **faker** - Fake data generation

## ## Project Test Structure

...

```
tests/
├── conftest.py # Shared fixtures
└── unit/ # Unit tests
 ├── conftest.py
 ├── test_runner_manager.py
 ├── test_gpu_detector.py
 └── test_fleeting_plugin.py
└── integration/ # Integration tests
 ├── conftest.py
 ├── test_docker_integration.py
 └── test_kubernetes_integration.py
└── e2e/ # End-to-end tests
```

```
| └── conftest.py
| └── test_full_pipeline.py
|── fixtures/ # Test data and fixtures
| ├── configs/
| ├── data/
| └── mocks/
└── utils/ # Test utilities
 ├── factories.py
 └── helpers.py
```
## Writing Unit Tests

### Basic Test Structure

```python
"""
Test suite for RunnerManager.

Documentation: docs/runners/README.md
"""

import pytest
from unittest.mock import Mock, patch, call
from autogit.runners import RunnerManager, ProvisionError

class TestRunnerManager:
 """
 Test suite for RunnerManager class.

 @pytest.fixture
 def docker_client(self):
 """Mock Docker client."""
 return Mock()

 @pytest.fixture
 def config(self):
 """Test configuration."""
 return {
 "max_instances": 10,
 "idle_timeout": 300,
 "architectures": ["amd64", "arm64"]
 }

 @pytest.fixture
 def manager(self, docker_client, config):
 """RunnerManager instance."""
 return RunnerManager(docker_client, config)

 def test_provision_amd64_runner(self, manager, docker_client):
 """Test provisioning an amd64 runner without GPU."""
 # Arrange
 expected_id = "runner-123"
 docker_client.containers.run.return_value.id = expected_id

 # Act
 runner_id = manager.provision("amd64")

 # Assert
 assert runner_id == expected_id
```

```

```

docker_client.containers.run.assert_called_once()
call_args = docker_client.containers.run.call_args
assert call_args[1]["image"].endswith(":amd64")

def test_provision_with_nvidia_gpu(self, manager, docker_client):
    """Test provisioning runner with NVIDIA GPU."""
    # Arrange
    expected_id = "runner-gpu-123"
    docker_client.containers.run.return_value.id = expected_id

    # Act
    runner_id = manager.provision("amd64", gpu_type="nvidia")

    # Assert
    assert runner_id == expected_id
    call_args = docker_client.containers.run.call_args[1]
    assert "device_requests" in call_args
    assert call_args["device_requests"][0].driver == "nvidia"

def test_provision_unsupported_architecture(self, manager):
    """Test provisioning with unsupported architecture raises error."""
    # Act & Assert
    with pytest.raises(ProvisionError, match="Unsupported architecture"):
        manager.provision("mips")

def test_provision_max_instances_exceeded(self, manager):
    """Test provisioning fails when max instances reached."""
    # Arrange
    manager._current_instances = manager._config["max_instances"]

    # Act & Assert
    with pytest.raises(ProvisionError, match="Max instances reached"):
        manager.provision("amd64")

@pytest.mark.asyncio
async def test_provision_async(self, manager):
    """Test async provisioning."""
    # Arrange
    runner_id = await manager.provision_async("amd64")

    # Assert
    assert runner_id is not None
```

```

#### ### Testing Best Practices

1. \*\*Use descriptive test names\*\*:
 

```

```python
# Good
def test_provision_fails_when_docker_unavailable(self):

# Bad
def test_provision_error(self):
```

```
2. \*\*Follow Arrange-Act-Assert pattern\*\*:
 

```

```python
def test_something(self):
    # Arrange - Set up test data
```

```

```

config = {"key": "value"}

Act - Execute the code under test
result = function_under_test(config)

Assert - Verify the results
assert result == expected
```

3. **One assertion per test** (when possible):
```python
Good - focused test
def test_provision_returns_valid_id(self):
 runner_id = manager.provision("amd64")
 assert runner_id.startswith("runner-")

Good - related assertions
def test_provision_configures_runner_correctly(self):
 runner_id = manager.provision("amd64")
 runner = manager.get_runner(runner_id)
 assert runner.architecture == "amd64"
 assert runner.status == "running"
```

4. **Use fixtures for setup**:
```python
@pytest.fixture
def sample_config():
 """Reusable configuration fixture."""
 return {
 "max_instances": 5,
 "idle_timeout": 300
 }
```

5. **Mock external dependencies**:
```python
@patch('autogit.runners.docker_client')
def test_with_mock(self, mock_docker):
 mock_docker.containers.run.return_value.id = "test-id"
 # Test implementation
```

### Testing Exceptions

```python
def test_provision_raises_error_on_invalid_config(self):
 """Test that invalid configuration raises ConfigError."""
 with pytest.raises(ConfigError) as exc_info:
 RunnerManager({})

 assert "max_instances" in str(exc_info.value)

def test_provision_logs_warning_on_gpu_unavailable(self, caplog):
 """Test warning logged when GPU unavailable."""
 with caplog.at_level(logging.WARNING):
 manager.provision("amd64", gpu_type="nvidia")

 assert "GPU not available" in caplog.text
```

```

```
```python
@ pytest.mark.asyncio
async def test_async_provision(self):
 """Test asynchronous runner provisioning."""
 runner_id = await manager.provision_async("amd64")
 assert runner_id is not None

@ pytest.mark.asyncio
async def test_concurrent_provisioning(self):
 """Test multiple concurrent provisions."""
 tasks = [
 manager.provision_async("amd64")
 for _ in range(5)
]
 runner_ids = await asyncio.gather(*tasks)
 assert len(runner_ids) == 5
 assert len(set(runner_ids)) == 5 # All unique
```

```

Writing Integration Tests

Integration tests verify that components work together correctly.

```
```python
"""Integration tests for Runner Manager with Docker.
"""

Documentation: docs/runners/README.md
"""

```

```
import pytest
import docker
from autogit.runners import RunnerManager
```

```
@pytest.fixture(scope="module")
def docker_client():
 """Real Docker client for integration testing."""
 return docker.from_env()
```

```
@pytest.fixture(scope="module")
def runner_manager(docker_client):
 """Runner manager with real Docker client."""
 config = {
 "max_instances": 3,
 "idle_timeout": 60,
 "image": "gitlab/gitlab-runner:alpine"
 }
 return RunnerManager(docker_client, config)
```

```
class TestDockerIntegration:
 """Integration tests with Docker."""

 def test_provision_real_runner(self, runner_manager):
 pass
```

```

"""Test provisioning a real Docker container."""
Act
runner_id = runner_manager.provision("amd64")

Assert
assert runner_id is not None

Verify container exists
container = runner_manager._docker.containers.get(runner_id)
assert container.status == "running"

Cleanup
runner_manager.deprovision(runner_id)

def test_provision_with_volumes(self, runner_manager):
 """Test runner provisioning with volume mounts."""
 # Act
 runner_id = runner_manager.provision(
 "amd64",
 volumes={"/cache": {"bind": "/cache", "mode": "rw"}}
)

 # Assert
 container = runner_manager._docker.containers.get(runner_id)
 mounts = container.attrs['Mounts']
 assert any(m['Destination'] == '/cache' for m in mounts)

 # Cleanup
 runner_manager.deprovision(runner_id)

def test_autoscaling_behavior(self, runner_manager):
 """Test runner autoscaling up and down."""
 # Arrange - simulate job queue
 job_queue = [
 {"arch": "amd64"},
 {"arch": "amd64"},
 {"arch": "arm64"}
]

 # Act - scale up
 runner_ids = []
 for job in job_queue:
 runner_id = runner_manager.scale_for_job(job)
 runner_ids.append(runner_id)

 # Assert - runners provisioned
 assert len(runner_ids) == 3

 # Act - scale down after idle timeout
 import time
 time.sleep(65) # Wait for idle timeout
 runner_manager.cleanup_idle()

 # Assert - runners deprovisioned
 active = runner_manager.get_active_runners()
 assert len(active) == 0
```
### Database Integration Tests

```

```
```python
@pytest.fixture
def database():
 """Test database fixture."""
 db = create_test_database()
 yield db
 db.drop_all()

def test_runner_state_persistence(runner_manager, database):
 """Test that runner state is persisted correctly."""
 # Arrange
 runner_id = runner_manager.provision("amd64")

 # Act
 runner_manager.save_state(database)
 runner_manager_new = RunnerManager.load_state(database)

 # Assert
 runners = runner_manager_new.get_active_runners()
 assert runner_id in [r.id for r in runners]
```

```

Writing End-to-End Tests

E2E tests verify complete user workflows.

```
```python
"""End-to-end tests for AutoGit platform.

Documentation: docs/tutorials/quickstart.md
"""

import pytest
import requests
from time import sleep

@pytest.fixture(scope="module")
def autogit_instance():
 """Deploy a complete AutoGit instance for testing."""
 # Start all services
 subprocess.run(["docker", "compose", "-f", "compose/test/docker-compose.yml", "up", "-d"])

 # Wait for services to be ready
 wait_for_gitlab()
 wait_for_traefik()

 yield "https://gitlab.test.local"

 # Cleanup
 subprocess.run(["docker", "compose", "-f", "compose/test/docker-compose.yml", "down", "-v"])

def wait_for_gitlab(timeout=300):
 """Wait for GitLab to be ready."""
 start = time.time()
 while time.time() - start < timeout:
```

```
try:
 response = requests.get("https://gitlab.test.local/-/health")
 if response.status_code == 200:
 return
except requests.exceptions.RequestException:
 pass
sleep(5)
raise TimeoutError("GitLab did not become ready")

class TestE2EWorkflow:
 """End-to-end workflow tests."""

def test_complete_cicd_pipeline(self, autogit_instance):
 """Test complete CI/CD pipeline from push to deployment."""
 # Arrange - Create project
 project = create_test_project(autogit_instance)

 # Act - Push code with .gitlab-ci.yml
 push_test_code(project)

 # Assert - Pipeline runs successfully
 pipeline = wait_for_pipeline(project)
 assert pipeline.status == "success"

 # Assert - Runner was auto-provisioned
 runners = get_active_runners()
 assert len(runners) > 0

 # Assert - Runner auto-deprovisioned after idle
 sleep(70)
 runners = get_active_runners()
 assert len(runners) == 0

def test_multi_arch_build_pipeline(self, autogit_instance):
 """Test multi-architecture build pipeline."""
 # Arrange
 project = create_test_project(autogit_instance)

 # Act - Push multi-arch build config
 push_multiarch_config(project)

 # Assert - Pipeline uses correct runners
 pipeline = wait_for_pipeline(project)
 jobs = pipeline.jobs

 amd64_job = next(j for j in jobs if "amd64" in j.name)
 arm64_job = next(j for j in jobs if "arm64" in j.name)

 assert amd64_job.runner.tags == ["amd64"]
 assert arm64_job.runner.tags == ["arm64"]

def test_gpu_accelerated_job(self, autogit_instance):
 """Test GPU-accelerated job scheduling."""
 # Arrange
 project = create_test_project(autogit_instance)

 # Act - Push GPU job config
 push_gpu_job_config(project)
```

```

Assert - Job runs on GPU-enabled runner
pipeline = wait_for_pipeline(project)
gpu_job = next(j for j in pipeline.jobs if "gpu" in j.name)

assert "nvidia" in gpu_job.runner.tags or "amd" in gpu_job.runner.tags
assert gpu_job.status == "success"

def test_sso_authentication(self, autogit_instance):
 """Test SSO authentication flow."""
 # Act - Attempt to access GitLab
 response = requests.get(f"{autogit_instance}/projects", allow_redirects=False)

 # Assert - Redirected to Authelia
 assert response.status_code == 302
 assert "authelia" in response.headers["Location"]

 # Act - Login via Authelia
 session = login_via_authelia("testuser", "testpass")

 # Assert - Can access GitLab
 response = session.get(f"{autogit_instance}/projects")
 assert response.status_code == 200
```
## Test Coverage

### Measuring Coverage

```bash
Run tests with coverage
pytest --cov=src --cov-report=html --cov-report=term

View HTML report
open htmlcov/index.html

Check coverage threshold
pytest --cov=src --cov-fail-under=80
```

### Coverage Configuration

```ini
[run]
source = src/
omit =
 /tests/
 /venv/
 /__pycache__/
 /site-packages/

[report]
exclude_lines =
 pragma: no cover
 def __repr__
 raise AssertionError
 raise NotImplementedError
 if __name__ == '__main__':
```

```

```
if TYPE_CHECKING:  
    @abstractmethod  
...  
  
### Coverage Goals  
  
- **Overall**: 80%+ coverage  
- **Critical paths**: 95%+ coverage (runner management, GPU detection)  
- **New code**: 90%+ coverage (enforced in CI)  
  
## Continuous Integration  
  
### GitHub Actions Integration  
  
```yaml  
.github/workflows/test.yml
name: Tests

on: [push, pull_request]

jobs:
 test:
 runs-on: ubuntu-latest
 strategy:
 matrix:
 python-version: ['3.11', '3.12']

 steps:
 - uses: actions/checkout@v4

 - name: Set up Python
 uses: actions/setup-python@v5
 with:
 python-version: ${{ matrix.python-version }}

 - name: Install dependencies
 run: |
 pip install uv
 uv sync

 - name: Run unit tests
 run: uv run pytest tests/unit/ --cov --cov-report=xml

 - name: Upload coverage
 uses: codecov/codecov-action@v3
 with:
 file: ./coverage.xml

 - name: Run integration tests
 run: uv run pytest tests/integration/

 - name: Check coverage threshold
 run: uv run pytest --cov --cov-fail-under=80
...

Test Data Management

Using Factories
```

```
```python
# tests/utils/factories.py
import factory
from autogit.models import Runner, Job

class RunnerFactory(factory.Factory):
    """Factory for creating test Runner instances."""

    class Meta:
        model = Runner

    id = factory.Sequence(lambda n: f"runner-{n}")
    architecture = "amd64"
    status = "running"
    gpu_type = None
    tags = factory.LazyAttribute(lambda obj: [obj.architecture])

class JobFactory(factory.Factory):
    """Factory for creating test Job instances."""

    class Meta:
        model = Job

    id = factory.Sequence(lambda n: n)
    name = factory.Sequence(lambda n: f"job-{n}")
    stage = "build"
    architecture = "amd64"
    requires_gpu = False

# Usage in tests
def test_runner_assignment():
    runner = RunnerFactory()
    job = JobFactory(architecture="amd64")

    assigned = assign_job_to_runner(job, runner)
    assert assigned is True
```

Using Faker for Realistic Data

```python
from faker import Faker
fake = Faker()

def test_with_realistic_data():
    """Test with realistic fake data."""
    config = {
        "project_name": fake.company(),
        "repo_url": fake.url(),
        "owner_email": fake.email()
    }

    project = create_project(config)
    assert project.name == config["project_name"]
```

```

```
Performance Testing
```

```
Load Testing
```

```
```python
import pytest
from locust import HttpUser, task, between

class GitLabUser(HttpUser):
    """Load test user for GitLab."""

    wait_time = between(1, 5)

    @task(3)
    def browse_projects(self):
        self.client.get("/projects")

    @task(1)
    def trigger_pipeline(self):
        self.client.post("/projects/1/pipeline", json={
            "ref": "main"
        })
```

```
# Run with: locust -f tests/performance/test_load.py --host=https://gitlab.test.local
```

```
```
```

```
Benchmarking
```

```
```python
def test_provision_performance(benchmark):
    """Benchmark runner provisioning speed."""
    result = benchmark(manager.provision, "amd64")
    assert result is not None
```

```
def test_gpu_detection_performance(benchmark):
    """Benchmark GPU detection speed."""
    result = benchmark(gpu_detector.detect_all)
    assert len(result) > 0
``
```

```
## Mocking Strategies
```

```
### Mocking Docker Client
```

```
```python
@pytest.fixture
def mock_docker():
 """Mock Docker client."""
 with patch('docker.from_env') as mock:
 client = Mock()
 mock.return_value = client

 # Configure mock behavior
 client.containers.run.return_value.id = "test-container-id"
 client.containers.list.return_value = []
```

```
 yield client
...

Mocking Kubernetes API

```python
@pytest.fixture
def mock_k8s_client():
    """Mock Kubernetes client."""
    with patch('kubernetes.client.CoreV1Api') as mock:
        api = Mock()
        mock.return_value = api

        # Configure mock behavior
        api.create_namespaced_pod.return_value = Mock(
            metadata=Mock(name="test-pod")
        )

    yield api
...

```

Mocking External APIs

```
```python
@pytest.fixture
def mock_gitlab_api():
 """Mock GitLab API responses."""
 with requests_mock.Mocker() as m:
 # Mock project endpoint
 m.get(
 'https://gitlab.test.local/api/v4/projects/1',
 json={'id': 1, 'name': 'test-project'}
)

 # Mock pipeline endpoint
 m.post(
 'https://gitlab.test.local/api/v4/projects/1/pipeline',
 json={'id': 123, 'status': 'pending'}
)

 yield m
...

```

#### ## Debugging Tests

##### ### Using pytest Debugging

```
```bash
# Drop into debugger on failure
pytest --pdb

# Drop into debugger at start of test
pytest --trace

# Show local variables on failure
pytest --showlocals

# Run specific test with verbose output

```

```
pytest -vv tests/unit/test_runner_manager.py::TestRunnerManager::test_provision
```

```
...
```

Logging in Tests

```
```python
import logging

def test_with_logging(caplog):
 """Test with log capture."""
 with caplog.at_level(logging.DEBUG):
 manager.provision("amd64")

 # Assert log messages
 assert "Provisioning runner" in caplog.text
 assert any("amd64" in record.message for record in caplog.records)
```
```

```

### ### Test Markers

```
```python
# Mark slow tests
@pytest.mark.slow
def test_long_running_operation():
    pass

# Mark GPU tests
@pytest.mark.gpu
@pytest.mark.skipif(not has_gpu(), reason="GPU not available")
def test_gpu_detection():
    pass

# Mark integration tests
@pytest.mark.integration
def test_docker_integration():
    pass

# Run specific markers
# pytest -m "not slow" # Skip slow tests
# pytest -m gpu       # Only GPU tests
```
```

```

Testing Checklist

Before submitting a PR, ensure:

- [] All tests pass locally
- [] New code has tests (90%+ coverage)
- [] Tests follow naming conventions
- [] Tests are documented
- [] Integration tests added for new components
- [] E2E tests added for new user workflows
- [] Performance tests added if applicable
- [] Tests run in CI/CD
- [] Coverage threshold met (80%+)
- [] No flaky tests
- [] Test documentation updated

Common Testing Pitfalls

1. Flaky Tests

Problem: Tests that pass/fail randomly
Solution: Avoid timing dependencies, use proper mocking

```
```python
Bad - timing dependent
def test_async_operation():
 start_async_task()
 time.sleep(1) # Hope it's done
 assert is_complete()

Good - wait for condition
def test_async_operation():
 start_async_task()
 wait_until(lambda: is_complete(), timeout=5)
 assert is_complete()
```

```

2. Test Interdependence

Problem: Tests that depend on execution order

Solution: Make each test independent

```
```python
Bad - depends on previous test
def test_create_runner():
 global runner_id
 runner_id = manager.provision("amd64")

def test_delete_runner():
 manager.deprovision(runner_id) # Depends on previous test

Good - independent tests
def test_create_and_delete_runner():
 runner_id = manager.provision("amd64")
 manager.deprovision(runner_id)
```

```

3. Over-Mocking

Problem: Mocking too much defeats the purpose of testing

Solution: Mock only external dependencies

```
```python
Bad - mocking internal logic
@patch('autogit.runners.RunnerManager._validate_config')
def test_provision(mock_validate):
 mock_validate.return_value = True
 # Not testing real validation

Good - testing real logic
def test_provision_with_invalid_config():
 with pytest.raises(ConfigError):
 RunnerManager({"invalid": "config"})
```

```

Resources

- **pytest Documentation**: <https://docs.pytest.org/>
- **Testing Best Practices**: <https://testdriven.io/blog/testing-best-practices/>
- **Coverage.py**: <https://coverage.readthedocs.io/>
- **Factory Boy**: <https://factoryboy.readthedocs.io/>

Next Steps

- Review [Coding Standards](standards.md)
- Set up [Development Environment](setup.md)
- Read [CI/CD Guide](ci-cd.md)
- Check [Common Tasks](common-tasks.md)

Documentation Version: 1.0.0

Last Updated: YYYY-MM-DD

Related Docs: [Development Overview](README.md) | [Standards](standards.md) | [CI/CD](ci-cd.md)

...

docs/architecture/adr/README.md

```markdown

# Architecture Decision Records

This directory contains Architecture Decision Records (ADRs) for AutoGit.

### ## What is an ADR?

An Architecture Decision Record (ADR) is a document that captures an important architectural decision made along with its context and consequences.

### ## When to Create an ADR

Create an ADR when making decisions about:

- \*\*Technology Choices\*\*: Selecting frameworks, libraries, or tools
- \*\*Architectural Patterns\*\*: Choosing design patterns or architectural styles
- \*\*Infrastructure Decisions\*\*: Deployment strategies, scaling approaches
- \*\*Integration Approaches\*\*: How components communicate
- \*\*Data Management\*\*: Storage solutions, data flow
- \*\*Security\*\*: Authentication, authorization, encryption choices

### ## ADR Format

Each ADR follows this structure:

```markdown

ADR-XXX: [Title]

Status: [Proposed | Accepted | Deprecated | Superseded]

Date: YYYY-MM-DD

Deciders: [List of people involved]

Technical Story: [Link to issue/epic]

Context

[Describe the context and problem statement]

Decision Drivers

- [Driver 1]
- [Driver 2]
- [Driver 3]

Considered Options

- [Option 1]
- [Option 2]
- [Option 3]

Decision Outcome

Chosen option: [Option X]

Positive Consequences

- [Positive consequence 1]
- [Positive consequence 2]

Negative Consequences

- [Negative consequence 1]
- [Negative consequence 2]

Pros and Cons of the Options

[Option 1]

- **Good**: [Advantage]
- **Bad**: [Disadvantage]
- **Neutral**: [Consideration]

[Option 2]

- **Good**: [Advantage]
- **Bad**: [Disadvantage]

Links

- [Related ADR](XXX-related-decision.md)
- [External resource]
- ...

ADR Index

| ADR Title Status Date |
|---|
| ----- ----- ----- ----- |
| [001](001-traefik-vs-nginx.md) Traefik vs NGINX Ingress Accepted 2024-12-21 |
| [002](002-fleet-ingress-plugin.md) Custom Fleet Ingress Plugin Design Accepted 2024-12-21 |
| [003](003-multi-architecture.md) Multi-Architecture Strategy Accepted 2024-12-21 |
| [004](004-sso-solution.md) SSO Solution Selection Accepted 2024-12-21 |
| [005](005-dns-management.md) DNS Management Approach Accepted 2024-12-21 |
| [006](006-storage-architecture.md) Storage Architecture Accepted 2024-12-21 |

Creating a New ADR

1. **Copy the template**:

```
```bash
cp docs/architecture/adr/template.md docs/architecture/adr/XXX-your-title.md
````
```
2. **Number sequentially**: Use the next available number
3. **Fill in all sections**: Don't leave sections empty
4. **Update this index**: Add your ADR to the table above
5. **Link related ADRs**: Cross-reference related decisions
6. **Get review**: Have the team review before marking as "Accepted"

ADR Lifecycle

...

Proposed → Accepted → [Deprecated | Superseded]

...

- **Proposed**: Under discussion
- **Accepted**: Decision made and implemented
- **Deprecated**: No longer recommended but still in use
- **Superseded**: Replaced by a newer ADR

Modifying Existing ADRs

- **Never modify accepted ADRs**: Create a new ADR that supersedes it
- **Update status**: Mark old ADR as "Superseded by ADR-XXX"
- **Link new ADR**: Reference the superseded ADR

Best Practices

1. **Be concise**: ADRs should be quick to read
2. **Be specific**: Include concrete examples
3. **Include context**: Explain why, not just what
4. **Document alternatives**: Show what you didn't choose and why
5. **Link resources**: Include research and references
6. **Update promptly**: Write ADRs when decisions are made, not after

Resources

- [ADR GitHub Organization](<https://adr.github.io/>)
- [Documenting Architecture Decisions](<https://cognitect.com/blog/2011/11/15/documenting-architecture-decisions>)
- [ADR Tools](<https://github.com/npryce/adr-tools>)

****Last Updated****: 2024-12-21
****Maintainer****: Project Team

docs/architecture/adr/001-traefik-vs-nginx.md

```markdown

**\*\*Status\*\*:** Accepted  
**\*\*Date\*\*:** 2024-12-21  
**\*\*Deciders\*\*:** Project Team  
**\*\*Technical Story\*\*:** Core infrastructure selection

## ## Context

AutoGit requires an ingress controller to manage external access to services in the Kubernetes cluster. The ingress controller must:

- Route traffic to multiple services (GitLab, Authelia, DNS)
- Provide SSL/TLS termination
- Support Let's Encrypt automation
- Be MIT-license compatible
- Be lightweight and performant
- Have active community support

## ## Decision Drivers

- **License Compatibility**: Must be MIT or compatible (Apache 2.0, BSD)
- **Maintenance Status**: Active development and security updates
- **SSL Automation**: Native Let's Encrypt support
- **Resource Usage**: Low memory and CPU footprint
- **Ease of Configuration**: Simple, declarative configuration
- **Community Support**: Active community and good documentation
- **Cloud-Native**: Built for Kubernetes/Docker environments

## ## Considered Options

1. **Traefik** (MIT License)
2. **NGINX Ingress Controller** (Apache 2.0)
3. **HAProxy Ingress** (Apache 2.0)
4. **Envoy/Contour** (Apache 2.0)

## ## Decision Outcome

**Chosen option**: **Traefik**

## ### Rationale

1. **NGINX EOL**: NGINX Ingress Controller will be retired in March 2026 with no further security updates
2. **Native Let's Encrypt**: Traefik has built-in ACME support without additional controllers
3. **Dynamic Configuration**: Automatic service discovery and configuration
4. **MIT License**: Perfect compatibility with project license
5. **Modern Architecture**: Built specifically for cloud-native environments
6. **Better DX**: Simpler configuration than NGINX
7. **Active Development**: Regular releases and security patches

## ### Positive Consequences

- No need to manage separate cert-manager for basic SSL (though we still use it for advanced features)
- Automatic service discovery reduces configuration overhead
- Built-in dashboard for monitoring
- Native Docker and Kubernetes support
- Lower learning curve for contributors
- Future-proof (no EOL concerns)

### ### Negative Consequences

- Different from traditional NGINX (team familiarity)
- Smaller ecosystem than NGINX (though growing)
- Different configuration paradigm (labels vs ConfigMaps)

### ## Pros and Cons of the Options

#### ### Traefik

- **Good**: MIT licensed, active development, native Let's Encrypt, automatic service discovery
- **Good**: Built-in dashboard, Gateway API support
- **Good**: Lower resource usage than NGINX
- **Good**: Simpler configuration model
- **Bad**: Smaller ecosystem than NGINX
- **Bad**: Less mature than NGINX (though stable)
- **Neutral**: Different configuration paradigm

#### ### NGINX Ingress Controller

- **Good**: Very mature, large ecosystem
- **Good**: Familiar to most teams
- **Good**: Extensive documentation
- **Bad**: **Being retired March 2026** - no security updates after
- **Bad**: More complex configuration
- **Bad**: Requires separate cert-manager for SSL
- **Bad**: Manual service configuration

#### ### HAProxy Ingress

- **Good**: Very performant
- **Good**: Mature and stable
- **Bad**: More complex to configure
- **Bad**: Less Kubernetes-native
- **Bad**: Smaller community than Traefik/NGINX
- **Neutral**: Apache 2.0 license

#### ### Envoy/Contour

- **Good**: Very powerful and flexible
- **Good**: Used by Istio/service meshes
- **Bad**: Much more complex than needed
- **Bad**: Higher resource usage
- **Bad**: Steeper learning curve
- **Neutral**: Overkill for this use case

### ## Implementation Notes

#### ### Traefik Configuration

```
```yaml
# Traefik deployed via Helm
helm install traefik traefik/traefik \
  --namespace traefik \
  --set ports.web.redirectTo=websecure \
  --set certificatesResolvers.letsencrypt.acme.email=admin@example.com
```
```

#### ### Service Discovery Example

```
```yaml
# Services auto-discovered via annotations
apiVersion: v1
kind: Service
metadata:
  name: gitlab
  annotations:
    traefik.ingress.kubernetes.io/router.entrypoints: websecure
    traefik.ingress.kubernetes.io/router.tls.certresolver: letsencrypt
spec:
  ports:
    - port: 80
```

Links

- [Traefik Documentation](https://doc.traefik.io/traefik/)
- [NGINX Ingress EOL Announcement](https://github.com/kubernetes/ingress-nginx/issues/10870)
- [Related: ADR-006 SSL/TLS Configuration](006-storage-architecture.md)
- [Traefik GitHub](https://github.com/traefik/traefik)

Supersedes

None (initial decision)

Superseded By

None (current)

Last Updated: 2024-12-21

Status: Accepted and Implemented

```

```